

Ivan Evdokimov, Ph.D.

Manchester, UK | ivevdm@gmail.com | linkedin.com/in/ivan-evdokimov | github.com/ivan020

Dear Hiring Manager,

I'm writing to apply for the AI Engineer (NLP/ML) role on Meltwater's Data Enrichment team. I've spent the last several years building and deploying NLP and ML systems in production, and I'd like to bring that experience to a team focused on turning unstructured social data into something clients can act on in real time.

Transformer-based NLP and LLMs are central to my recent work. At the UK Data Service I trained, fine-tuned, and deployed PyTorch, scikit-learn, and LLaMA models via Hugging Face, vLLM, and Ollama for automated metadata classification, reaching 95%+ F1 accuracy on sensitive social science datasets. That work covered the full lifecycle your role touches — tokenization and classification through to fine-tuning and transfer learning on pretrained models — and it had to run reliably against real, messy text at scale rather than clean benchmark data.

I also have direct experience with retrieval-augmented generation and agentic LLM workflows. I built a self-hosted RAG pipeline (Chat-LLM) using LangChain, pgvector, and MiniLM embeddings to let users query past content conversationally, and at the UK Data Service I delivered a LangGraph-based agent-to-agent workflow with hallucination checks to automate exploratory data analysis. Both projects required the same judgment this role asks for: getting a model's outputs to be genuinely useful and trustworthy in a live product, not just accurate in isolation.

On the engineering side, I'm comfortable with the infrastructure this role expects. I've built ETL pipelines processing 100K+ record datasets, deployed ML retraining pipelines on AWS Lambda and Step Functions with drift monitoring, and worked with Dockerised deployments as part of an end-to-end MLOps setup. I haven't used Terraform or Kubernetes extensively, but I pick up infrastructure tooling quickly — my career has consistently involved learning the pipeline layer needed to get models into production, not just building the models themselves.

I'd welcome the chance to talk about how this background could contribute to Meltwater's Data Enrichment team. Thank you for your time and consideration.

Kind regards,
Ivan Evdokimov