

Ivan Evdokimov, Ph.D.

Manchester, UK | ivevdm@gmail.com | linkedin.com/in/ivan-evdokimov | github.com/ivan020

Dr Fernandez De Arroyabe Arranz, Dr Lee, Dr Fejos, and Mr Bodill,

I am writing to apply for the Technical Business Analyst – AI & Compliance (KTP Associate) role with the University of Essex and Pimberly Software Development Ltd (Ref: REQ10157).

I am a software engineer with a PhD in Computational Finance from the University of Essex. I build backend systems, cloud infrastructure, and machine learning pipelines.

In my previous role, I designed REST APIs with FastAPI and built backend systems with LLM integrations. I also ran CI/CD pipelines with GitHub Actions and Vercel, and deployed Dockerised MLOps infrastructure. This reduced deployment time to two days, added continuous production monitoring, and increased the quality of the generated code, improving maintainability and reducing technical debt.

In addition, I led UX design and user testing for a Java mobile app used by 20 field technicians. I gathered feedback from the technicians and used it to revise the interface and workflow. The app was successfully deployed and went on to save three weeks of manual data collection a year.

At the UK Data Service, I trained and deployed machine learning models using PyTorch, scikit-learn, and LLaMA. These models classified metadata and assessed disclosure risk, following ONS statistical disclosure control (SDC) methodology. I reached 95%+ F1 accuracy on sensitive social science datasets. I also migrated disclosure control algorithms from R to C++, achieving an 11x performance gain and enabling real-time processing of datasets with 500,000+ records.

Outside of work, I build personal projects to learn new technologies and see how they behave in production, not just in theory. Chat-LLM is a self-hosted RAG pipeline running on a Raspberry Pi. It uses LangChain, pgvector, and Ollama to let Twitch viewers query past stream transcripts in real time. This allowed me to learn how RAG systems behave in production and how to optimise complex pipelines to reduce latency.

FinBot is a chat-based trading simulation. It has a Golang backend, a PostgreSQL database, and a React frontend, and it has served over 120 monthly users through Twitch chat. This application was deployed to other Twitch streamers, allowing them to diversify their content and build community around financial topics. It also allows a non-finance audience to learn about stock investing without risking their own funds.

My most recent project is MTG Card Reader, a Golang REST API that reads Magic: The Gathering cards using OCR, identifies them with an LLM, and returns live price data through a React frontend. I am currently expanding the application to include a chatbot that fetches card data and gives recommendations on a card's suitability for a user's deck. This would also allow for more card discoverability: a user could input a desired outcome and get card recommendations.

I am interested in the chance to build the Digital Product Passport module for Pimberly, and to help prepare their platform for ESPR ahead of the regulation coming into force. I would welcome the opportunity to discuss my application further.

Kind regards,
Ivan Evdokimov